


Adversarial Examples

Deep Learning: Bryan Pardo & Patrick O'Reilly, Northwestern University, Spring 2022


Learned decision surfaces are
messy and don't align with
human intuition

Looks good, right?




Learned decision surface for XOR problem

Let's zoom out a little




Learned decision surface for XOR problem

Zooming out more...




What are these things really learning?


Spirals data



Decision surface a human might draw




Actual decision surface learned by a network



These decision surfaces that don't align with human decision surfaces make networks brittle & easy to fool


What if we “nudge” an example over the line?

- Gradient descent alters the decision boundary
- Adversarial attacks alter the input
- Do it right and a human won’t see a difference
- ...but the machine might really screw up a classification



In 2 dimensions, a bad surface is obvious


- What about in 2 million dimensions?



One of these was labeled “panda” by a trained net.
The other was labeled “bucket”. Which is which?

Image from: <https://www.borealisai.com/en/blog/advertorch-adversarial-training-tool-implement-attack-and-defence-strategies/>

The one on the right is a “perturbed” image



PANDA



BUCKET

Gradient Descent Pseudocode

Initialize $\theta^{(0)}$

Repeat until stopping condition met:

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta} L(X, Y; \theta^{(t)})$$

Return $\theta^{(t_{max})}$

$\theta^{(t)}$ are the parameters of the model at time step t

$\nabla_{\theta} L(X, Y; \theta^{(t)})$ is the gradient of the loss function with respect to model parameters $\theta^{(t)}$

η controls the step size

$\theta^{(t_{max})}$ is the set of parameters that did best on the loss function.

Just flip which thing we're optimizing

Initialize $X^{(0)}$

Repeat until stopping condition met:

$$X^{(t+1)} = X^{(t)} + \eta \nabla_X L(X^{(t)}, Y; \theta)$$

Return $X^{(enough)}$

$X^{(t)}$ is an example at time t

$\nabla_X L(X^{(t)}, Y; \theta)$ is the gradient of the loss function with respect to example features $X^{(t)}$

η controls the step size

$X^{(enough)}$ is the minimal change needed to flip the category of X

Even Simpler: Fast Gradient Sign method

$$X^{(t+1)} = X^{(t)} + \eta sign(\nabla_X L(X^{(t)}, Y; \theta))$$

EXPLAINING AND HARNESSING ADVERSARIAL EXAMPLES

Ian J. Goodfellow, Jonathon Shlens & Christian Szegedy, ICLR 2015

Gradient Sign attack


- The pixels are all independent dimensions
- Find the gradient in the pixel space
- Add (clipped) noise along the gradient (a little noise for every pixel)
- Do it right and the image looks the same to the user...
...but looks entirely different to the network.

That same thing in pictures

$$f(x + \delta_0) = \text{PANDA}$$

$$f(x + \delta_1) = \text{PANDA}$$

$$f(x + \delta_{\dots}) = \text{PANDA}$$




That same thing in pictures

$$f(x + \delta_0) = \text{PANDA}$$

$$f(x + \delta_1) = \text{PANDA}$$

$$f(x + \delta_{\dots}) = \text{PANDA}$$

$$f(x + \delta_N) = \text{GIBBON}$$



Yes, it's just that easy



x
“panda”
57.7% confidence

$$+ .007 \times$$



$\text{sign}(\nabla_x J(\theta, x, y))$
“nematode”
8.2% confidence

$$=$$



$x +$
 $\epsilon \text{sign}(\nabla_x J(\theta, x, y))$
“gibbon”
99.3 % confidence

(Image from Goodfellow et al 2014)

Why.....

-does this gradient-based attack make sense?
- ...did they use the sign of the gradient multiplied by a fixed step size, instead of the actual gradient?

Defending against attacks

Our problem...

- An attacker can straightforwardly force the classifier to recategorize inputs.
- This could be a big problem for...
 - Self driving cars
 - Verification of identity
 - Etc..
- What can we do about it?



$$f(\mathbf{x} + \delta_0) = \text{"Panda"}$$



$$f(\mathbf{x} + \delta_1) = \text{"Horse"}$$



$$f(\mathbf{x}) = \text{"Panda"}$$




$$f(\mathbf{x} + \delta_2) = \text{"Gibbon"}$$

Defenses

- Security through obscurity (don't let them see your weights)
 - Can be helpful....not a guarantee. There are black-box attacks.
- Randomly modify the input to screw up the perturbation.
 - At training time (use adversarial examples in training)
 - At inference time (we'll talk more about this)
- Ensembling
 - Train N different networks with different architectures & training data
 - Use majority voting for classification
 - Hope they can't attack a majority of them simultaneously

Force them into the open

- If we think that our attacks will be nudges to put images just over the border....
- ..and these nudges are designed to be imperceptible.
- Perhaps we can force them to make a perceptible change if they want to force misclassification...



Randomized smoothing

- At inference, random perturbations are sampled from a Gaussian centered at the input \mathbf{x} and a majority vote is taken from the classifier's predictions over these perturbed inputs.

